

PREDICTIVE ANALYTICS WITH ATHLETICS DATA

Week 1 · How this class works, and a statistics review

TUESDAY AUG 25 AND THURSDAY AUG 27, 2026



CASTLE SPORTS



INTRODUCE YOURSELF.

Name, major, and year.

Why you registered for this class.

The sport you follow most closely, and the job you would take tomorrow if somebody offered it.

01

HOW THIS CLASS WORKS

MY GOAL: EQUIP YOU WITH THE INTUITION, CRITICAL THINKING, TECHNICAL, AND COMMUNICATION SKILLS NEEDED TO USE ANALYTICS TO ADD VALUE TO AN ORGANIZATION.

EVERY UNIT HAS FOUR PARTS

1 Tuesday's lecture.

The unit's concepts, in class.

2 Thursday, in class, by hand, no agentic AI.

We will use Excel to do the calculations and build an understanding of our approaches. Without a mental model of what we're doing and why, you will misapply models to the wrong contexts and communicate their outputs incorrectly. You may still use an LLM to assist with troubleshooting formulas and understanding our approach.

3 The Case Study. On your own, with Claude Code.

You will practice using agentic coding tools to quickly run similar analysis more efficiently. When complete, you will communicate your findings in your own words, along with submitting the code.

4 Tuesday's quiz, closed-book.

A short quiz at the start of Tuesday's class to confirm that you understand the concepts.

ASSIGNMENT ASSISTANCE RESTRICTIONS

1 Thursday, in class, by hand, no agentic AI.

You may use an LLM to assist with troubleshooting formulas and understanding our approach. The hand-build may be finished at home. It is submitted by Monday at 11:59pm, along with the AI-assisted assignment and brief.

2 The Case Study. On your own, with Claude Code.

Not allowed: submitting analysis created or edited by another human. You may discuss assignments with classmates, and use books, instructors, and LLMs. Not allowed: submitting briefs written by AI or another human. The BRIEF.md submissions must be in your own words; copying and pasting from an LLM and changing some of the verbiage is also in violation. You need to be able to sit in a room with your boss, a coach, or a GM and explain your findings.

3 Tuesday's quiz, closed-book.

Quizzes and both exams are closed-book and proctored. If you can only explain a concept by consulting an outside resource, you don't understand the concept yet.

GRADING

COMPONENT	WEIGHT	NOTES
Quizzes	25%	Weekly, Tuesday, closed-book. Two lowest dropped.
Case Studies	25%	Equal halves: the analysis and the brief. Audit and quiz artifacts required.
Exams	25%	Exam 1 multiple-choice (Oct 8). Exam 2 is a live audit (Nov 19).
Final project	25%	Proposal Week 11, presented Dec 8. No final exam.

FOUR COMPONENTS OF EACH CASE STUDY

1 The analysis.

The spreadsheet, code, chart, table. Whatever you built, including Thursday's hand-built file.

2 The quiz transcript.

`/quiz-me` quizzes you on that week's concepts and writes the whole exchange into your repo. This will prepare you for the in-class quiz on Tuesday.

3 The audit record.

`/audit` reviews your analysis like a skeptical senior analyst and writes down each finding. You respond to its critiques either by making a change or defending your methodology, in the same file.

4 The brief.

`BRIEF.md`, half a page to the week's audience, typed by you.

HOW A CASE STUDY IS GRADED

1 Half: the analysis.

The spreadsheet, code, chart, table. Whatever you built, including Thursday's committed hand-built file.

2 Half: BRIEF.md, typed by you.

Half a page to the week's named audience, usually a coach. Graded on decision quality and on honesty about what you do not know. You may ask Claude to red-team the draft; you may not ask it to write the draft.

3 Required but not graded on content: the audit record and quiz transcript.

They catch issues with your analysis before I do and prepare you for the quiz. Nothing blocks the push, so you can submit without them if you are out of time, but a missing one counts against the analysis half of the grade.

THE FIRST HALF OF THE SEMESTER

WK	TUE / THU	CONTENT
1	Aug 25 / 27	Onboarding and statistics review
2	Sep 1 / 3	Classical probability and inference
3	Sep 8 / 10	Bayesian statistics
4	Sep 15 / 17	Descriptive and probabilistic statistics
5	Sep 22 / 24	Introduction to regression
6	Sep 29 / Oct 1	Regression continued
7	Oct 6 / 8	Review Tuesday. Exam 1 Thursday.

First Case Study due Mon Aug 31. Labor Day pushes Week 2's to Tue Sep 8. Sep 8 is the last add or drop day.

THE SECOND HALF OF THE SEMESTER

WK	TUE / THU	CONTENT
8	Oct 13 / 15	Data collection and responsible automation
9	Oct 20 / 22	Storytelling, visualization, skills and hooks
10	Oct 27 / 29	Logical fallacies and AI shortcomings
11	Thu Nov 5 only	Monte Carlo. Final project proposal due.
12	Nov 10 / 12	Machine learning
13	Nov 17 / 19	Practical exam: three flawed analyses, audited live
14	Dec 1 / 3	Final project work sessions
15	Tue Dec 8	Final presentations

Election Day suspends classes Tue Nov 3. No class Nov 23 to 27 for Thanksgiving. Final deliverable due Mon Dec 14.

02

OUR TECH STACK

THREE THINGS BEFORE YOU LEAVE TODAY

- 1** Claude Pro, installed and signed in.
\$20/month for four months. If that creates an issue, talk to or email me and we'll get it covered.
- 2** Claude Code and Git on your machine.
We'll walk through the setup page together today.
- 3** Your repository, created from the course template on GitHub.
This is where every assignment lives for the rest of the semester. Your copy is yours; I can see it, your classmates cannot.

TWO WORDS YOU WILL USE EVERY WEEK: COMMIT AND PUSH

Commit: a checkpoint.

It snapshots your work with a label, like saving your game, and you can always come back to it.

Push: upload your commits to GitHub.

Until you push, your work exists only on your laptop and cannot be seen or graded.

The rhythm is always the same.

Save the file, commit the checkpoint, push it to GitHub.

FOUR CUSTOM COMMANDS FOR THIS COURSE

/teach

Explains a concept using the week's data. Adapted from Matt Pocock's /teach skill.

/audit

Has the agent audit your own work like a skeptical senior analyst, re-deriving every key number from the raw data. Run this before you push your assignments.

/quiz-me

Quizzes you on that week's concepts before you submit your assignment. This will prepare you for the in-class quiz on Tuesday.

/submit

Submits your work. Before doing so, it checks every required file and tells you if anything's missing.

SETUP TIME.

- 1 Install Git and Python.
- 2 Install the Claude Desktop app. Sign in.
- 3 GitHub account. Create your repo from the template. Add jackdavl.
- 4 Install the Python packages.
- 5 Fill in About me in MISSION.md. Commit and push it.

EVERY TUESDAY STARTS THE SAME WAY: PULL THE NEW COURSE MATERIALS

- 1 Open a terminal and go to the course-materials folder.

```
cd spax402/spax402-course-materials
```

- 2 Pull the week's new files.

```
git pull
```

- 3 That is where every deck and data file arrives.

Your own repo holds your work; this one holds the handouts. You cloned it just now, so your first pull is next Tuesday.

03

THE SHAPE OF THE DATA


Mean, median, mode, and the reason the first question about any statistic is what its distribution looks like.

MEAN, MEDIAN, MODE

The average, the middle value, and the most common value.
In a symmetric distribution they land on the same number;
how far apart they sit is a measure of how skewed the
data is.


Mean = $\sum x / n$. Median = the 50th percentile. Mode = the most frequent value, and the only one of the three that also works on categories.

WHERE THE THREE-SIGMA RULES COME FROM



Symmetric, one peak, mean equals median equals mode. Roughly 68, 95, and 99.7 percent of values inside one, two, and three standard deviations. Every one of those numbers is a property of this exact shape.

POINTS PER TEAM-GAME IS CLOSE TO NORMAL



Mean 23.01, median 23, skew 0.30 across 570 team-games. Close enough to normal that the sigma rules roughly hold. Many sport statistics are not.

nflverse play-by-play, 2024 regular season and playoffs

? RUSHING YARDS PER CARRY: DO YOU EXPECT THE MEAN TO BE HIGHER OR LOWER THAN THE MEDIAN?

SAQUON BARKLEY'S MEDIAN CARRY WAS THE SAME IN 2024 AND 2025

	2024	2025
Carries	436	306
Yards per carry	5.74	4.07
Median carry	3 yards	3 yards
Gained 4 or more yards	48.2%	45.1%
Runs of 15 or more yards	32	14
Yards from his best 10% of carries	49.8%	45.1%

Two different backs by the mean, 1.67 yards apart, with an identical median carry. Roughly half his yards come from a tenth of his carries in both seasons, so the mean is reporting 32 explosive runs becoming 14.


nflverse play-by-play, regular season and playoffs. Rushing attempts only.
SPAX 402 · Week 1

WHICH OF THESE STATISTICS DO YOU EXPECT TO BE SKEWED, IN WHICH DIRECTION, AND WHY?

Home runs. Free-throw percentage. Golf scores. Time to recover from an injury.
Height. Salary.

For each one: is there a floor, is there a ceiling, and what would it take for a single athlete to sit far out on the tail?

TWO PEAKS USUALLY MEAN YOU HAVE MIXED TWO DIFFERENT THINGS TOGETHER



NBA field-goal percentage, every qualifying player. One peak in the mid-40s, a second at 50 to 52. A distribution with two modes is bimodal, and the second mode is a question that needs an answer.

Severini (2020), Analytic Methods in Sports

? WHAT DO YOU THINK THAT SECOND PEAK OF PLAYERS ON THE RIGHT IS?

A BIMODAL DISTRIBUTION MEANS YOUR DATA LIKELY HAS MULTIPLE UNDERLYING SUBCLASSES.

Find the subclass before you compute anything else. Every summary statistic you calculate before you split is an average of two populations that do not belong in the same average.

04

HOW TO MEASURE SPREAD

STANDARD DEVIATION

Roughly the average distance of a data point from the mean. Small means the values cluster; large means they scatter.

The square root of the mean squared distance from the mean. Squaring is what stops the positive and negative distances from cancelling, and it is also why one extreme value moves an SD more than it moves a median.

COEFFICIENT OF VARIATION

The standard deviation divided by the mean, so spread can be compared across statistics measured on completely different scales.

$CV = SD / \text{mean}$. Unitless, so it survives a change of units.



**WHAT LEAGUE HAS THE MOST VARIATION IN SCORING BETWEEN
TEAMS?**

MLB, NFL, NBA, OR NHL?

THE RAW NUMBERS: SCORING PER GAME, LEAGUE BY LEAGUE

LEAGUE	STATISTIC	MEAN	SD
MLB	Runs scored	4.3	0.508
NFL	Points scored	22.2	5.49
NBA	Points scored	98.5	4.04
NHL	Goals scored	2.8	0.268

Rank these leagues by raw SD and decide whether that ranking makes sense, given how differently each league scores.

05

CALCULATING THE SPREAD OF QUALITATIVE VARIABLES

The variation in a pitcher's pitch deployment, a coach's play calling, or a defense's formations matters. We need a way to measure it.

YOU CANNOT TAKE THE STANDARD DEVIATION OF “RUN, PASS, PASS, RUN.”

There is nothing to average. But a play-caller who runs on every short-yardage snap is obviously more predictable than one who splits it fifty-fifty, so the variation is real and it needs a number.

ENTROPY MEASURES HOW SPREAD OUT A SET OF SHARES IS

$H = - \sum p \cdot \ln(p)$ over every category


Max $H = \ln(n)$ the value when every share is equal

Standardized entropy = $H / \ln(n)$

In plain English: p is one category's share of the total (a run rate of 60% is $p = 0.60$), $\ln$ is the natural log (=LN() in Excel), $\sum$ means add up one $p \cdot \ln(p)$ term per category, and n is how many categories there are. A category with a share of zero contributes zero, not an error. Do not try to take the log of nothing.

RANGES 0 TO 1: 0 IS FULLY PREDICTABLE, 1 IS AN EVEN SPLIT.

KERSHAW THREW ONE PITCH 62.5% OF THE TIME



Two more pitches, the sinker and the cutter, are in the arsenal and were thrown zero times in this sample. Whether they count toward n is a decision you make and then state.

Published Statcast pitch-mix shares

? ON THE ZERO-TO-ONE SCALE, WHERE DOES THIS PITCH MIX LAND?



COULD A TEAM'S RUN/PASS SPLIT ENTROPY COME OUT AT EXACTLY ZERO? WHAT WOULD THAT LOOK LIKE?

06

THIS WEEK'S WORK

THURSDAY, IN CLASS: THE HAND-BUILD. THE DATA IS IN YOUR REPO AT WEEKS/WEEK01/DATA.

1 NFL entropy: one team, in Excel, from the formula.

In your own repo, open `weeks/week01/data/hand-build-nfl-entropy-worksheet.xlsx` in Excel. It holds every Philadelphia Eagles run or pass play on third or fourth down with exactly two yards to go, from the 2024 regular season and playoffs. That is twenty-three plays. Fill in every highlighted cell to compute the standardized entropy of that run/pass mix from the formula.

2 Save the Excel file in `weeks/week01/` of your repo, then commit and push.

Later, when the agent builds the same numbers from the full data, it has to match your hand-built answer before you accept anything it builds.

THE EXCEL FUNCTIONS YOU NEED FOR THE HAND-BUILD

- 1** **COUNTIFS** counts how many plays fall in each category.
Point it at the `play_type` column and ask for “run”, then ask again for “pass”. Your two counts should add up to twenty-three.
- 2** **SUM** adds your terms together, which is the summation step in the formula.
Use it for the total play count, and again to add the $p \cdot \ln(p)$ terms. The minus sign goes in front of that sum.
- 3** **LN** gives you the natural log.
 $\ln(\text{number})$ is the $\ln$ in the formula. Use it on each share, and on the number of categories to get the maximum you divide by.

THE CASE STUDY, STEP BY STEP

1 The question.

Do less predictable play-callers have more success in short yardage situations?

2 Open Claude Code in your repo folder, spax402-your-username.

Start here, not in a terminal. Spec the job in your own words: the question, the target, the assumptions. Claude Code will send you to a terminal if it needs to. (/teach is optional, if you want the week's concepts walked through first.)

3 The data.

A script to pull the 2024 play-by-play is already in your repo at weeks/week01/data/. Have your agent run it, then ask what fields the data holds and decide which ones your question needs.

THE CASE STUDY, STEP BY STEP (CONTINUED)

4 Direct the agent.

Run the pull script, filter to third and fourth and short, and compute each team's run/pass split, standardized entropy, and whatever you decided success means. Tell it to ask you any clarifying questions before it starts, and answer them yourself.

5 Verify against Thursday.

Put your hand-build Excel file into `weeks/week01/` in your repo. The agent must reproduce your team and match your number before you accept the league table.

6 Ship it.

Save the team table as `.xlsx` in `outputs/`, plus one chart. Write `BRIEF.md` in your own words. Run `/audit` and `/quiz-me`, then `/submit`, which checks and pushes. Upload the repository's link to Canvas by Monday at 11:59pm.

WHAT GOES IN THE BRIEF.

How did you define success in short yardage, and why that measure?

Do you see a trend between standardized entropy and your measure of success?

What other information would you want to look at to determine how important entropy is for success?

How would you go about assessing the optimal entropy for a given team?

Half a page, in BRIEF.md, due Monday at 11:59pm.